

# Homework 3: Diffusion Models

CSC 8851 – Deep Learning - Fall 2025

## Learning Objectives

- Implement a linear noise schedule and visualize the forward diffusion process.
- Use **diffusers** components: **UNet2DModel**, **DDPMScheduler**, and **DDIMScheduler**.
- Train a class-conditional denoiser that predicts  $x_0$  and implement classifier-free guidance (CFG).
- Sample with DDPM (full steps) and DDIM (short trajectory) using CFG; compare speed/quality.

## Assignment Tasks (100 pts)

### Part 1: Linear Noise Schedule and Forward Noising (10 pts)

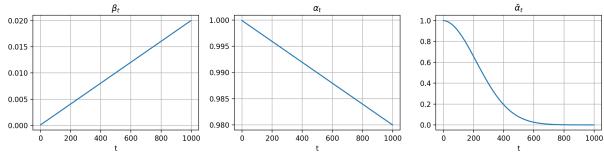
Let  $T = 1000$ . Define per-step quantities

$$\beta_t \in (0, 1), \quad \alpha_t = 1 - \beta_t, \quad \bar{\alpha}_t = \prod_{i=0}^t \alpha_i, \quad t = 0, \dots, T-1,$$

with a **linear** progression  $\beta_t \in [10^{-4}, 2 \cdot 10^{-2}]$ . The forward (noising) process is

$$q(\mathbf{x}_t \mid \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}) \iff \mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \sim \mathcal{N}(0, \mathbf{I}).$$

- (5 pts) Compute and plot  $\{\beta_t\}$ ,  $\{\alpha_t\}$ , and  $\{\bar{\alpha}_t\}$  vs.  $t$  (three plots). See Figure 1a.
- (5 pts) For one MNIST image, visualize  $\mathbf{x}_t$  at `steps=[0, 50, 100, 200, 250, 300, 400, 500, 600, 700, 800, 999]`. See Figure 1b. **Note:** Scale images to  $[-1, 1]$ .



(a) Linear schedule:  $\beta_t$ ,  $\alpha_t$ , and  $\bar{\alpha}_t$ .



(b) Forward noising grid over the specified steps.

Figure 1: Part 1 diagnostics.

*Note (not required):* A *cosine* schedule designs  $\bar{\alpha}_t$  via a smooth target curve that preserves more signal early and can stabilize training; we only implement **linear** here.

### Part 2: Model and Schedulers (PyTorch Diffusers) (20 pts)

- **UNet2DModel**: convolutional U-Net that accepts an input tensor and a discrete timestep  $t$ . In our setup it predicts  $\hat{\mathbf{x}}_0$  (`prediction_type="sample"`).

**UNet2DModel configuration (Do not change the architecture).**

```
sample_size = 32,  in_channels = 1 + 10,  out_channels = 1
layers_per_block = 2,  block_out_channels = (32, 64, 128),
down_block_types = ("DownBlock2D", "DownBlock2D", "DownBlock2D")
up_block_types = ("UpBlock2D", "UpBlock2D", "UpBlock2D").
```

**UNet2DModel architecture.** Input  $\mathbf{Z} \in \mathbb{R}^{B \times 11 \times 32 \times 32}$  flows through:

1. **Time embedding:** sinusoidal  $\rightarrow$  MLP  $\rightarrow$  injected into ResNet blocks.
2. **Encoder (Down path):** three stages with two ResNet blocks each; channels  $11 \rightarrow 32 \rightarrow 64 \rightarrow 128$ ; spatial downsample  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$ . Skip connections saved.
3. **Mid block:** ResNet(s) at  $(B, 128, 4, 4)$ .
4. **Decoder (Up path):** three stages with skip concatenations; channels  $128 \rightarrow 64 \rightarrow 32$ ; spatial upsample  $4 \rightarrow 8 \rightarrow 16 \rightarrow 32$ .
5. **Output head:**  $3 \times 3$  conv  $\rightarrow (B, 1, 32, 32)$  giving  $\hat{\mathbf{x}}_0$  (no final activation).

- **DDPMScheduler:**

- Training (forward noising): `Xt = ddpm.add_noise(X0, eps, t)` implements  $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon$ .
- Sampling (reverse): `ddpm.set_timesteps(N)` then iterate  $t \in ddpm.timesteps$  and call `ddpm.step(model_output=hat_x_0, timestep=t, sample=X)` to obtain the next sample.
- With `prediction_type="sample"`, `model_output` must be  $\hat{\mathbf{x}}_0$  (your network output).

- **DDIMScheduler:**

- Same API as DDPM but supports a shortened, non-Markovian trajectory; set `ddim.set_timesteps(K)` with  $K \ll T$ .
- The call `ddim.step(model_output=hat_x_0, timestep=t, sample=X, eta=0.0)` yields the deterministic DDIM update.

### DDPMScheduler and DDIMScheduler configuration

```
num_train_timesteps = T,  beta_schedule = "linear",  prediction_type = "sample"
```

- **Class conditioning (one-hot planes).** Given labels  $y \in \{0, \dots, 9\}^B$ , build one-hot class maps  $\mathbf{C} \in \mathbb{R}^{B \times 10 \times H \times W}$  and concatenate with the (noised) image  $\mathbf{X}_t \in \mathbb{R}^{B \times 1 \times H \times W}$  to form the model input:

$$\mathbf{Z}_t = [\mathbf{X}_t; \mathbf{C}] \in \mathbb{R}^{B \times 11 \times H \times W}, \quad H = W = 32.$$

*Efficient construction:* `F.one_hot(y, 10).float().view(B, 10, 1, 1).expand(B, 10, H, W)` (no per-pixel loops).

`make_model_input` (**training-time CFG dropout**). We randomly zero all class channels (dropout) for some samples to train the model to also handle *unconditional* inputs:

$$\text{make\_model\_input}(\mathbf{X}_t, y, p_{\text{uncond}}) = [\mathbf{X}_t ; \mathbf{C} \odot \mathbf{M}], \quad \mathbf{M} \in \{0, 1\}^{B \times 1 \times 1 \times 1}, \Pr[M=0] = p_{\text{uncond}}.$$

```
def make_model_input(x, y, p_uncond=0.1):
    B, _, H, W = x.shape
    C = F.one_hot(y, 10).float().view(B, 10, 1, 1).expand(B, 10, H, W)
    m = (torch.rand(B, device=x.device) < p_uncond).float().view(B, 1, 1, 1)
    C = C * (1.0 - m) # zero all class planes for some samples
    return torch.cat([x, C.to(x.dtype)], dim=1) # (B, 11, H, W)
```

### How to use it.

- **Training:** always feed `make_model_input(Xt, y, p_uncond)` into `UNet2DModel`, where `Xt` are noised images generated by `DDPMScheduler.add_noise`. The model predicts  $\hat{\mathbf{x}}_0$ ; optimize  $\text{MSE}(\hat{\mathbf{x}}_0, \mathbf{x}_0)$ .
- **Sampling with CFG (Part 5):** build two inputs at each step: *conditional*  $[\mathbf{X}_t; \mathbf{C}]$  and *unconditional*  $[\mathbf{X}_t; \mathbf{0}]$ , run the model on both, and combine

$$\hat{\mathbf{x}}_0^{\text{cfg}} = \hat{\mathbf{x}}_0^{\text{cond}} + s(\hat{\mathbf{x}}_0^{\text{cond}} - \hat{\mathbf{x}}_0^{\text{uncond}}), \quad s \geq 0.$$

### Grading for Part 2 (20 pts).

- (10 pts) Correct instantiation of `UNet2DModel`, `DDPMScheduler`, and `DDIMScheduler` with the specified configuration; Evaluate a dummy forward: `model(Z, t).sample`  $\rightarrow$   $(B, 1, 32, 32)$  and verify shape.
- (10 pts) Implement `make_model_input` (correct shapes and dropout behavior).

### Part 3: Training (50 epochs) and Curve (20 pts)

Train the UNet to predict  $\mathbf{x}_0$ :

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\varepsilon}, \quad \mathcal{L} = \mathbb{E}[\|\hat{\mathbf{x}}_0(\mathbf{x}_t, y, t) - \mathbf{x}_0\|_2^2].$$

### One-epoch training pseudocode ( $\mathbf{x}_0$ prediction).

```
# Assumes: model (UNet2DModel), ddpm (DDPMScheduler), train_loader, optimizer 'opt'
model.train()
running, total, count = None, 0.0, 0
for X0, y in train_loader: # X0 in [-1,1], shape (B,1,32,32)
    X0 = X0.to(dev); y = y.to(dev).long()
    B = X0.size(0)

    # Sample t and construct forward noising (DDPM)
    t = torch.randint(0, ddpm.config.num_train_timesteps, (B,), device=dev, dtype=torch.long)
```

```

eps = torch.randn_like(X0)
Xt = ddpm.add_noise(X0, eps, t)
# where x_t = sqrt(bar_alpha_t) x0 + sqrt(1-bar_alpha_t) eps

# Build conditional/unconditional-batch input with CFG dropout
Zt = make_model_input(Xt, y, p_uncond=0.1)      # (B,11,32,32)

# Predict x0 and optimize MSE(x0_hat, x0)
x0_hat = model(Zt, t).sample
loss = F.mse_loss(x0_hat, X0)

opt.zero_grad(set_to_none=True)
loss.backward()
opt.step()

# (Optional) track loss
total += float(loss.item()) * B; count += B
running = loss.item() if running is None else 0.98*running + 0.02*loss.item()

avg_loss = total / max(count, 1)
print(f"avg_epoch_loss={avg_loss:.4f}, running={running:.4f}")

```

- (15 pts) Implement loop: sample  $t$ , form  $\mathbf{x}_t$  via `ddpm.add_noise`, build inputs with `make_model_input`, predict  $\hat{\mathbf{x}}_0$ , optimize MSE.
- (5 pts) Plot the training loss. See Figure 2.

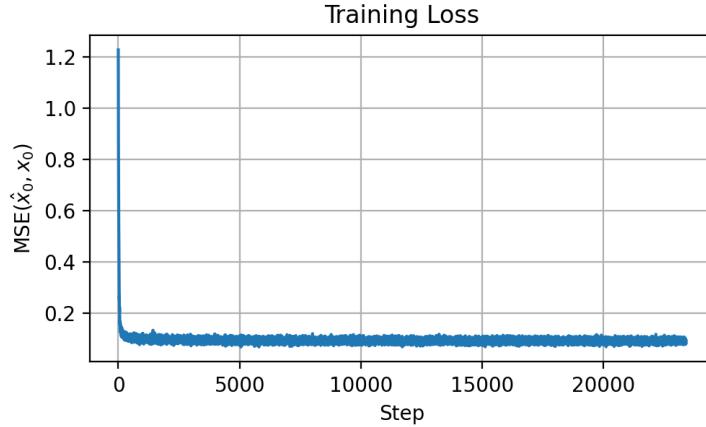


Figure 2: Training loss curve (MSE).

#### Part 4: Visualizing $\hat{x}_0$ Across Noise (20 pts)

Fix one  $(\mathbf{x}_0, y)$ . For the timesteps  $[0, 50, 100, 200, 250, 300, 400, 500, 600, 700, 800, 999]$ , form  $\mathbf{x}_t$  using the forward process and visualize the model's prediction  $\hat{\mathbf{x}}_0(\mathbf{x}_t, \cdot, t)$  in two

settings:

- (10 pts) **Conditional**. Use the *conditional* input (true label; no dropout): set  $p_{\text{uncond}}=0$  inside your visualization helper so the class maps are active. See Figure: 3a.
- (10 pts) **Unconditional**. Use the *unconditional* input by concatenating zero class maps (i.e.,  $[\mathbf{x}_t; \mathbf{0}]$ ). See Figure: 3b.



(a) Conditional  $\hat{x}_0$  across timesteps (top:  $t=0$ ; rightmost:  $t=999$ ).



(b) Unconditional  $\hat{x}_0$  across the same timesteps (zero class maps).

Figure 3: Part 5:  $\hat{x}_0$  predictions across noise levels (conditional vs. unconditional).

### Part 5: Sampling — DDPM and DDIM with CFG (30 pts)

At each step, form cond/uncond inputs, get  $\hat{\mathbf{x}}_0^{\text{cond}}$ ,  $\hat{\mathbf{x}}_0^{\text{uncond}}$ , then  $\hat{\mathbf{x}}_0^{\text{cfg}}$  with scale  $s$ .

- **5.1 DDPM Sampling** (15 pts). Use `DDPMScheduler` (`prediction_type="sample"`) and its reverse-time indices to iterate and update `X <- ddpm.step(...).prev_sample`. Generate a  $10 \times 8$  class grid. See Figure 4a.
- **5.2 DDIM Sampling** (15 pts). DDIM uses a shortened, often deterministic ( $\eta=0$ ) trajectory and works with a DDPM-trained model. Use 50 steps,  $\eta=0$ , and the same CFG combination; generate the same class grid and compare speed/quality vs. DDPM. See Figure 4b.

## Submission Instructions

1. Submit a Jupyter Notebook named `CSC8851_F2025_HW3_<YourName>.ipynb` with all code and plots (runs top-to-bottom).
2. Submit a PDF export named `CSC8851_F2025_HW3_<YourName>.pdf`. Do not paste code screenshots. **Make sure your PDF version has the figures.**
3. Upload both files (`.ipynb` and `.pdf`) to iCollege.

*You may include additional qualitative results (e.g., failure cases). Cite any external sources used.*

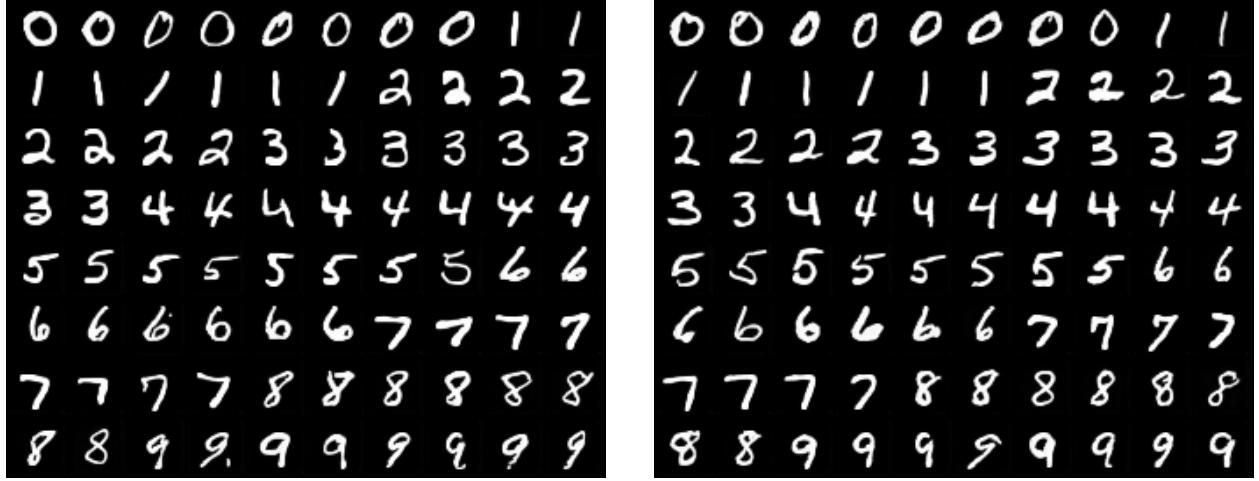
(a) DDPM + CFG ( $s = 2.0$ ).(b) DDIM( $50, \eta=0$ ) + CFG ( $s = 2.0$ ).

Figure 4: Conditional sampling results (rows: digits 0–9; 8 samples each).

## Appendix: DDIM Sampling Details and Speed

**Setup.** Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=0}^t \alpha_i$ . Given a denoiser that predicts either  $\hat{\mathbf{x}}_0(\mathbf{x}_t, t)$  or noise  $\hat{\boldsymbol{\varepsilon}}(\mathbf{x}_t, t)$ , we can convert between the two via

$$\hat{\boldsymbol{\varepsilon}}(\mathbf{x}_t, t) = \frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \hat{\mathbf{x}}_0}{\sqrt{1 - \bar{\alpha}_t}}, \quad \hat{\mathbf{x}}_0(\mathbf{x}_t, t) = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \hat{\boldsymbol{\varepsilon}}}{\sqrt{\bar{\alpha}_t}}. \quad (1)$$

### DDPM vs. DDIM (one-step updates)

**DDPM (Markovian, stochastic).** A common parameterization of the reverse step is

$$\mathbf{x}_{t-1}^{(\text{DDPM})} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \hat{\boldsymbol{\varepsilon}}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}, \quad \mathbf{z} \sim \mathcal{N}(0, \mathbf{I}), \quad (2)$$

where  $\sigma_t^2$  is chosen from the DDPM posterior variance family. This update is *Markovian* and typically run at all  $T$  steps (e.g.,  $T=1000$ ).

**DDIM (non-Markovian, implicit).** DDIM defines a family of trajectories that share the same single-step marginals as DDPM but allow skipping steps and optionally removing noise. For a chosen subsequence of indices

$$\mathcal{S} = \{\tau_1 > \tau_2 > \dots > \tau_K\}, \quad K \ll T,$$

the DDIM step from  $\tau_k$  to  $\tau_{k+1}$  is

$$\sigma_{\tau_k}^{(\text{DDIM})} = \eta \sqrt{\frac{1 - \bar{\alpha}_{\tau_{k+1}}}{1 - \bar{\alpha}_{\tau_k}}} \sqrt{1 - \frac{\bar{\alpha}_{\tau_k}}{\bar{\alpha}_{\tau_{k+1}}}}, \quad (3)$$

$$\mathbf{x}_{\tau_{k+1}}^{(\text{DDIM})} = \sqrt{\bar{\alpha}_{\tau_{k+1}}} \hat{\mathbf{x}}_0 + \sqrt{1 - \bar{\alpha}_{\tau_{k+1}} - (\sigma_{\tau_k}^{(\text{DDIM})})^2} \hat{\boldsymbol{\varepsilon}} + \sigma_{\tau_k}^{(\text{DDIM})} \mathbf{z}, \quad \mathbf{z} \sim \mathcal{N}(0, \mathbf{I}), \quad (4)$$

with  $\hat{\mathbf{x}}_0$  and  $\hat{\boldsymbol{\varepsilon}}$  evaluated at  $(\mathbf{x}_{\tau_k}, \tau_k)$  and related by (1). The hyperparameter  $\eta \in [0, 1]$  controls stochasticity.

**Deterministic DDIM (fast default).** Setting  $\eta = 0$  yields

$$\mathbf{x}_{\tau_{k+1}} = \sqrt{\bar{\alpha}_{\tau_{k+1}}} \hat{\mathbf{x}}_0 + \sqrt{1 - \bar{\alpha}_{\tau_{k+1}}} \hat{\boldsymbol{\varepsilon}} \quad (\eta = 0). \quad (5)$$

No per-step noise is injected; the trajectory becomes a deterministic mapping guided by the denoiser.

### Where the speed-up comes from

- **Fewer model calls.** DDPM typically uses all  $T$  small steps. DDIM chooses a *coarse* timestep set  $\mathcal{S}$  (e.g.,  $K=50$ ) and updates using (4). Complexity drops from  $\mathcal{O}(T)$  to  $\mathcal{O}(K)$  denoiser evaluations (e.g.,  $\sim 20\times$  fewer if  $K=50$ ,  $T=1000$ ).
- **Stable large steps.** Because the DDIM update directly re-targets the next marginal with the pair  $(\bar{\alpha}_{\tau_k}, \bar{\alpha}_{\tau_{k+1}})$  and the predicted  $(\hat{\mathbf{x}}_0, \hat{\boldsymbol{\varepsilon}})$ , it tolerates larger jumps in  $t$  than the stochastic DDPM posterior step.
- **Optional determinism.** With  $\eta=0$  (Eq. (5)), trajectories are noise-free and invertible in the ideal case, which improves consistency at small  $K$  and further reduces sampling variance.

### DDPM Generation (Pseudocode)

Algorithm: DDPM Sampling (x0-prediction) with optional CFG

Inputs:

- Trained model that predicts  $x_0$ :  $x0\_hat = model(Z, t).sample$
- Labels  $y$  ( $B, 1$ ), image shape  $(B, 1, H, W)$  with  $H=W=32$  for MNIST
- Full timestep set  $S_{full} = \{T-1, \dots, 0\}$  (usually  $T=1000$ )
- Guidance scale  $s \geq 0$  (use  $s=0$  to disable CFG)
- Precomputed  $\alpha[t]$ ,  $\alpha_{bar}[t]$ , and posterior variance  $\tilde{\beta}[t]$

Initialize:

```
X <- Normal(0, I)           # X has shape (B, 1, H, W)

for each t in S_full:        # reverse order: T-1, T-2, ..., 0

    # ----- (A) Predict x0 (conditional + unconditional for CFG) -----
    Zc <- concat(X, class_maps(y))           # (B, 11, H, W)
    x0_cond <- model(Zc, t).sample          # \hat{x}_0^{cond}

    Zu <- concat(X, zeros_like(class_maps(y))) # unconditional input
    x0_uncond <- model(Zu, t).sample          # \hat{x}_0^{uncond}

    x0_hat <- x0_cond + s * (x0_cond - x0_uncond) # CFG combine (s=0 disables)

    # ----- (B) Convert \hat{x}_0 -> predicted noise \hat{\epsilon} (optional) -----
    eps_hat <- (X - sqrt(alpha_bar[t]) * x0_hat) / sqrt(1 - alpha_bar[t])
```

```

# ----- (C) DDPM update (two equivalent forms) -----

# Form 1: noise-parameterized mean + stochasticity
mu = (1 / sqrt(alpha[t])) * (X - (1 - alpha[t]) / sqrt(1 - alpha_bar[t]) * eps_hat)
sigma = sqrt(tilde_beta[t])                                # posterior std (DDPM original choice)
X_prev = mu + sigma * Normal(0, I)

# Form 2: x0/epsilon mixture (algebraically equivalent)
# c = sqrt(1 - alpha_bar[t-1] - tilde_beta[t])    # define c for convenience
# X_prev = sqrt(alpha_bar[t-1]) * x0_hat + c * eps_hat
#           + sqrt(tilde_beta[t]) * Normal(0, I)

X <- X_prev

# Output (denormalize as needed)
return X

```

### Remarks.

- DDPM is *Markovian* and injects noise at every step via the posterior variance  $\tilde{\beta}_t$  (or a scaled variant).
- In practice, we rarely hand-code  $\mu$  and  $\sigma$ ; the scheduler computes the reverse step given  $\hat{x}_0$  (or  $\hat{\epsilon}$ ) and the current sample.

### DDIM Generation (Pseudocode)

Algorithm: DDIM Sampling (x0-prediction) with optional CFG

#### Inputs:

- Trained model that predicts  $x_0$ :  $x0\_hat = \text{model}(Z, t).\text{sample}$
- Labels  $y$  ( $B, \cdot$ ), image shape  $(B, 1, H, W)$  with  $H=W=32$  for MNIST
- Short timestep set  $S = \{\tau_1 > \tau_2 > \dots > \tau_K\}$ ,  $K \ll T$  (e.g.,  $K=50$ )
- Guidance scale  $s \geq 0$  (use  $s=0$  to disable CFG)
- DDIM stochasticity  $\eta$  in  $[0, 1]$  (use  $\eta=0$  for deterministic DDIM)
- Precomputed  $\alpha_{bar}[t]$  for all  $t$  (cumulative product of alphas)

#### Initialize:

```

X <- Normal(0, I)          # X has shape (B, 1, H, W)

for k = 1 to K-1:
    t      <- tau_k
    t_prev <- tau_{k+1}

# ----- (A) Predict x0 (conditional + unconditional for CFG) -----

```

```

Zc <- concat(X, class_maps(y))                      # (B, 11, H, W)
x0_cond <- model(Zc, t).sample                     # \hat{x}_0^{cond}

Zu <- concat(X, zeros_like(class_maps(y)))          # unconditional input
x0_uncond <- model(Zu, t).sample                   # \hat{x}_0^{uncond}

x0_hat <- x0_cond + s * (x0_cond - x0_uncond)    # CFG combine (s=0 disables)

# ----- (B) Convert \hat{x}_0 to predicted noise \hat{\epsilon} -----
eps_hat <- (X - sqrt(alpha_bar[t]) * x0_hat) / sqrt(1 - alpha_bar[t])

# ----- (C) DDIM variance and update -----
# sigma controls stochasticity of the step (DDIM general form):
sigma = eta
    * sqrt( (1 - alpha_bar[t_prev]) / (1 - alpha_bar[t]) )
    * sqrt( 1 - alpha_bar[t] / alpha_bar[t_prev] )

c = sqrt( 1 - alpha_bar[t_prev] - sigma^2 )

# DDIM step from t -> t_prev:
X_prev = sqrt(alpha_bar[t_prev]) * x0_hat
    + c * eps_hat
    + sigma * Normal(0, I)                         # drop this term if eta=0

X <- X_prev

# Output (denormalize for viewing as needed):
return X

```

## Using diffusers: DDPM vs. DDIM (x<sub>0</sub>-prediction)

DDPM (full steps).

```

# Assume: model predicts x0 (prediction_type="sample")
ddpm.set_timesteps(ddpm.config.num_train_timesteps)    # typically T=1000
X = Normal(0, I)                                       # (B,1,32,32)
for t in ddpm.timesteps:
    Zc = concat(X, class_maps(y))
    Zu = concat(X, zeros_like(class_maps(y)))
    x0_c = model(Zc, t).sample
    x0_u = model(Zu, t).sample
    x0_cfg = x0_c + s * (x0_c - x0_u)                  # CFG combine
    step = ddpm.step(model_output=x0_cfg, timestep=t, sample=X)
    X = step.prev_sample
# X is the generated image batch

```

**DDIM** (short trajectory, optional determinism).

```
# Use far fewer steps (e.g., K=50) and optionally set eta=0 for deterministic DDIM
ddim.set_timesteps(K)                                     # K << T, e.g., 50
X = Normal(0, I)                                         # (B,1,32,32)
for t in ddim.timesteps:
    Zc = concat(X, class_maps(y))
    Zu = concat(X, zeros_like(class_maps(y)))
    x0_c = model(Zc, t).sample
    x0_u = model(Zu, t).sample
    x0_cfg = x0_c + s * (x0_c - x0_u)                      # CFG combine
    step = ddim.step(model_output=x0_cfg, timestep=t, sample=X, eta=0.0)
    X = step.prev_sample
# X is the generated image batch
```

### Key differences.

- **DDPM:** Markovian, stochastic reverse process, typically run for all  $T$  steps; higher compute cost.
- **DDIM:** Non-Markovian implicit update; can jump across a subset of timesteps  $K \ll T$  and be deterministic ( $\eta=0$ ), giving a large speed-up with comparable quality.

### Notes.

- **Difference vs. DDPM.** DDPM uses a Markovian posterior update with per-step noise and is typically run for all  $T$  steps. DDIM uses a *non-Markovian* implicit update that can jump across a *subset* of steps  $\mathcal{S}$  (e.g., 50) and can be *deterministic* for  $\eta = 0$ .
- **Speed-up.** Reducing the number of denoiser calls from  $T$  to  $K$  yields  $\mathcal{O}(T/K)$  speed-up (e.g.,  $1000 \rightarrow 50$  gives  $\sim 20\times$  fewer steps) with competitive quality.
- **Disabling CFG.** Set  $s = 0$  to use unconditional sampling (or provide only the conditional branch and set  $x_0^{\text{uncond}} = 0$ ).
- **Shapes.** For MNIST:  $X \in \mathbb{R}^{B \times 1 \times 32 \times 32}$ , class maps  $(B, 10, 32, 32)$ , model input  $Z \in \mathbb{R}^{B \times 11 \times 32 \times 32}$ .